

LAB1

A.1 Prerequisites for working with Kubernetes

Install minikube on linux virtual fedora:

Requirement:

<https://minikube.sigs.k8s.io/docs/drivers/none/>

Installation guide:

<https://minikube.sigs.k8s.io/docs/start/>

WARNING:

Use minikube as **root** user

```
minikube start --force --driver=docker
```

Set kubectl cli alias

<https://minikube.sigs.k8s.io/docs/handbook/kubectl/>

```
cp /etc/bashrc /etc/bashrc.bk
```

```
echo "alias kubectl=\"minikube kubectl --\" >> /etc/bashrc
```

```
source /etc/bashrc
```

Start kubernetes in docker with NON user root

as root:

```
adduser developer
```

Create a password for developer

```
passwd developer
```

Add developer user to wheel and docker groups

```
usermod -aG wheel developer
```

```
usermod -aG docker developer
```

Login as developer

```
su - developer
```

Start minikube

```
minikube start --driver=docker
```

Set kubectl cli alias

<https://minikube.sigs.k8s.io/docs/handbook/kubectl/>

```
echo "alias kubectl=\"minikube kubectl --\" >> ~/.bashrc
```

Install kubernetes in you own operating system only if necessary

<https://minikube.sigs.k8s.io/docs/start/>

LAB2

Working with pod

CREATING A NAMESPACE FROM A YAML FILE

First, create a **custom-namespace.yaml** file with the following listing's contents (you'll find the file in the book's code archive).

```
apiVersion: v1
kind: Namespace
metadata:
  name: custom-namespace
```

Creating a simple YAML descriptor for a pod nginx-pod.yaml

```
kind: Pod
apiVersion: v1
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.11
    ports:
    - containerPort: 80
      protocol: TCP
```

Using `kubectl` create to create the pod

To create the pod from your YAML file, use the `kubectl create` command:

```
$ kubectl create -f nginx-pod.yaml -n custom-namespace
```

Use the following command to see the full descriptor of the pod and take a little time to inspect the output

```
$ kubectl get po nginx -o yaml -n custom-namespace
```

Let's list pods to see their statuses:

```
$ kubectl get po -n custom-namespace
NAME                READY    STATUS    RESTARTS    AGE
nginx                1/1     Running   0            2m14s
```

Set up the following alias:

```
cp /etc/bashrc /etc/bashrc.bk
echo "alias kcd=\"kubectl config set-context \$(kubectl config
current-context) --namespace \"" >> /etc/bashrc
source /etc/bashrc
```

Set the custom-namespace as default

```
$ kcd custom-namespace
Context "minikube" modified.
$ kubectl get po
NAME    READY    STATUS    RESTARTS    AGE
nginx   1/1     Running   0            84s
```

FORWARDING A LOCAL NETWORK PORT TO A PORT IN THE POD

The following command will forward your machine's local port 8080 to port 80 of your nginx pod:

```
$ kubectl get po
NAME                READY    STATUS    RESTARTS    AGE
nginx                1/1     Running   0            171m
$ kubectl port-forward nginx 8080:80
Forwarding from 127.0.0.1:8080 -> 80
Forwarding from [::1]:8080 -> 80
```

Open a new shell to get the log

```
$ kubectl logs po/nginx -f
```

The port forwarder is running and you can now connect to your pod through the local port.

```
$ curl localhost:8080
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
  }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
```

<p>If you see this page, the nginx web server is successfully installed and working. Further configuration is required.</p>

<p>For online documentation and support please refer to
nginx.org.

Commercial support is available at
nginx.com.</p>

<p>Thank you for using nginx.</p>
</body>
</html>

See the log:

```
$ kubectl logs po/nginx -f
127.0.0.1 - - [29/Nov/2020:21:46:19 +0000] "GET / HTTP/1.1" 200
425 "-" "Mozilla/5.0 (Macintosh; Intel Mac OS X 11_0_0)
AppleWebKit/537.36 (KHTML, like Gecko) Chrome/86.0.4240.198
Safari/537.36" "-"
127.0.0.1 - -
[...]
```

Deleting pods by deleting the whole namespace

```
$ kcd default
Context "minikube" modified.
```

Stop the port forwarder typing Ctrl+c

Delete your namespace "custom-namespace"

```
$ kubectl delete ns custom-namespace
namespace "custom-namespace" deleted
```

```
$ kubectl get ns
NAME                STATUS    AGE
default             Active   57d
kube-node-lease     Active   57d
kube-public         Active   57d
kube-system         Active   57d
```

LAB3

Creating a ReplicationController

create a YAML file called **kubia-rc.yaml** for your ReplicationController

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: kubia
spec:
  replicas: 3
  selector:
    app: kubia
  template:
    metadata:
      labels:
        app: kubia
    spec:
      containers:
      - name: kubia
        image: luksa/kubia
        ports:
        - containerPort: 8080
```

Use the `kubectl create` command to create your ReplicationController

```
$ kubectl create -f kubia-rc.yaml
replicationcontroller/kubia created
$ kubectl get po
```

NAME	READY	STATUS	RESTARTS	AGE
kubia-9cbh4	1/1	Running	0	116s
kubia-d7c8m	1/1	Running	0	116s
kubia-hwmt8	1/1	Running	0	116s

Try to delete a pod to see how the ReplicationController spins up a new one immediately

```
$ kubectl delete pod kubia-9cbh4
pod "kubia-9cbh4" deleted
```

```
$ kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
kubia-9cbh4	1/1	Terminating	0	13m
kubia-9gjdj	1/1	Running	0	27s
kubia-d7c8m	1/1	Running	0	13m
kubia-hwmt8	1/1	Running	0	13m

Now, let's see what information the `kubectl get` command shows for ReplicationControllers

```
$ kubectl get rc
NAME          DESIRED    CURRENT    READY    AGE
kubia         3          3          3        15m
```

You can see additional information about your ReplicationController with the `kubectl describe` command, as shown in the following listing.

```
$ kubectl describe rc kubia
Name:          kubia
Namespace:     default
Selector:      app=kubia
Labels:        app=kubia
Annotations:   <none>
Replicas:      3 current / 3 desired
Pods Status:   3 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  app=kubia
  Containers:
    kubia:
      Image:      luksa/kubia
      Port:       8080/TCP
      Host Port:  0/TCP
      Environment: <none>
      Mounts:      <none>
  Volumes:      <none>
Events:
  Type      Reason              Age   From              Message
  ----      -
  Normal    SuccessfulCreate    17m   replication-controller Created
pod: kubia-hwmt8
  Normal    SuccessfulCreate    17m   replication-controller Created
pod: kubia-9cbh4
  Normal    SuccessfulCreate    17m   replication-controller Created
pod: kubia-d7c8m
  Normal    SuccessfulCreate    5m1s  replication-controller Created
pod: kubia-9gjdj
```

Your ReplicationController has been keeping three instances of your pod running. Try to scale that number up to 10 now manually.

```
$ kubectl scale rc kubia --replicas=10
```

```
$ kubectl get po
NAME          READY    STATUS              RESTARTS    AGE
kubia-9gjdj    1/1     Running             0           8m45s
kubia-bdt9h    0/1     ContainerCreating   0           7s
```

kubia-bmz6g	0/1	ContainerCreating	0	7s
kubia-d7c8m	1/1	Running	0	21m
kubia-hwmt8	1/1	Running	0	21m
kubia-pbm7g	1/1	Running	0	7s
kubia-pp2lc	0/1	ContainerCreating	0	7s
kubia-tvt92	1/1	Running	0	7s
kubia-wzxng	0/1	ContainerCreating	0	7s
kubia-zgh9r	0/1	ContainerCreating	0	7s

Scaling your ReplicationController by editing it's definition

```
$ kubectl edit rc kubia
replicationcontroller/kubia edited
```

Set the value of

Spec:

Replicas: 3

```
$ kubectl get po
```

NAME	READY	STATUS	RESTARTS	AGE
kubia-9gjdj	1/1	Running	0	14m
kubia-d7c8m	1/1	Running	0	27m
kubia-hwmt8	1/1	Running	0	27m

Delete your ReplicationController without deleting associated pod (--cascade=false)

```
$ kubectl delete rc kubia --cascade=false
replicationcontroller "kubia" deleted
```

```
$ kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
kubia-9gjdj	1/1	Running	0	18m
kubia-d7c8m	1/1	Running	0	30m
kubia-hwmt8	1/1	Running	0	30m

You'll rewrite your ReplicationController into a ReplicaSet by creating a new file called **kubia-replicaset.yaml** with the contents in the following listing.

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: kubia
spec:
  replicas: 3
  selector:
    matchLabels:
      app: kubia
```

```

template:
  metadata:
    labels:
      app: kubia
  spec:
    containers:
      - name: kubia
        image: luksa/kubia

```

```

$ kubectl create -f kubia-replicaset.yaml
replicaset.apps/kubia created

```

```

$ kubectl describe rs kubia

```

```

Name:          kubia
Namespace:     default
Selector:      app=kubia
Labels:        <none>
Annotations:   <none>
Replicas:      3 current / 3 desired
Pods Status:   3 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  app=kubia
  Containers:
    kubia:
      Image:      luksa/kubia
      Port:       <none>
      Host Port:  <none>
      Environment: <none>
      Mounts:     <none>
  Volumes:      <none>
Events:        <none>

```

Try to delete an other pod:

```

$ kubectl delete pod kubia-9gjdj
pod "kubia-9gjdj" deleted

```

```

$ kubectl get po

```

NAME	READY	STATUS	RESTARTS	AGE
kubia-9gjdj	1/1	Terminating	0	24m
kubia-d7c8m	1/1	Running	0	37m
kubia-dv7h9	1/1	Running	0	5s
kubia-hwmt8	1/1	Running	0	37m

Clean all.

LAB 4

If not already active create a new ReplicaSet

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: kubia
spec:
  replicas: 3
  selector:
    matchLabels:
      app: kubia
  template:
    metadata:
      labels:
        app: kubia
    spec:
      containers:
        - name: kubia
          image: luksa/kubia
```

```
$ kubectl create -f kubia-replicaset.yaml
replicaset.apps/kubia created
```

Create a file called **kubia-svc.yaml** with the following listing's contents.

```
apiVersion: v1
kind: Service
metadata:
  name: kubia
spec:
  ports:
    - port: 80
      targetPort: 8080
  selector:
    app: kubia
```

After posting the YAML, you can list all Service resources in your namespace and see that an internal cluster IP has been assigned to your service:

```
$ kubectl get service
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubia	ClusterIP	10.102.158.76	<none>	80/TCP	16s

REMOTELY EXECUTING COMMANDS IN RUNNING CONTAINERS

be sure to replace the pod name and the service IP with your own:

```
$ kubectl get pod
NAME          READY   STATUS    RESTARTS   AGE
kubia-b2v5k   1/1     Running   0           57m
kubia-hst84   1/1     Running   0           57m
kubia-l55vf   1/1     Running   0           57m

$ kubectl exec kubia-hst84 -- curl -s http://10.102.158.76
You've hit kubia-l55vf
```

Using DNS

You can use the `kubectl exec` command to run `bash` (or any other shell) inside a pod's container.

```
$ kubectl get po
NAME          READY   STATUS    RESTARTS   AGE
kubia-b2v5k   1/1     Running   0           74m
kubia-hst84   1/1     Running   0           74m
kubia-l55vf   1/1     Running   0           74m

kubectl exec -it kubia-b2v5k bash
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in
a future version. Use kubectl exec [POD] -- [COMMAND] instead.
root@kubia-b2v5k:/# curl http://kubia.default.svc.cluster.local
You've hit kubia-l55vf
root@kubia-b2v5k:/# curl http://kubia.default
You've hit kubia-hst84
```

Look at the `/etc/resolv.conf` file in the container and you'll understand:

```
root@kubia-b2v5k:/# cat /etc/resolv.conf
nameserver 10.96.0.10
```

